

# FRAA: A Retrieval-Augmented Agent for Explainable Financial Risk Assessment

**Abstract.** Financial risk assessment depends on probabilities that are well calibrated, sensitive to recent external evidence, and sufficiently transparent for human review. In many deployed settings, risk models still rely on manually engineered tabular features and tree ensembles. These models remain effective for stable feature stores, but they have limited access to long behavioral histories and to market or regulatory information that changes after training. This paper presents the Financial Risk Assessment Agent (FRAA), a retrieval-augmented model for explainable financial risk assessment. For each scoring request, FRAA encodes heterogeneous user behavior with a time-aware Transformer, retrieves timestamp-valid documents from a financial knowledge base, fuses the retrieved evidence through cross-attention, and outputs both a calibrated risk score and a concise textual justification. On a proprietary 24-month multi-source dataset with over 20 million users, 1.2 billion interactions, and 15 data sources, FRAA achieves a Log Loss of 0.5163 and Recall@5 of 0.4185. These results correspond to a 19.6% Log Loss reduction relative to a deployed Random Forest baseline and a 9.4% reduction relative to a non-retrieval Transformer. Ablations show complementary contributions from dynamic behavior, temporal encoding, and timestamp-valid retrieval; feature-group occlusion identifies knowledge-retrieved context as the strongest feature-family signal. Retrieval-depth analysis further shows that five retrieved documents provide the best trade-off between coverage and noise. Further evaluations cover explanation quality rated 4.32/5 by domain experts and 15.7 ms per-query latency under the tested serving configuration.

**Keywords:** Financial risk assessment · Retrieval-augmented generation · Explainable AI · Risk scoring · Time-aware Transformer

## 1 Introduction

Financial institutions assess user-level risk from evidence that is heterogeneous, time-varying, and often incomplete at the moment of decision. Such evidence includes transaction histories, account states, macroeconomic indicators, regulatory disclosures, and news sentiment. Reliable risk estimates support credit review, fraud prevention, margin control, and compliance workflows [16,38,11,6,31]. In practice, many deployed systems still use carefully engineered tabular features with tree-based models such as Random Forests and gradient boosting [7,13,9,25]. These models remain strong operational baselines. However, their predictive state is usually fixed by the feature store, and they do not naturally incorporate long behavioral histories or newly released external information.

Neural sequence models and Transformer architectures offer a more flexible way to model user behavior over time [17,10,39,22,36,21]. Yet a purely parametric sequence model remains limited in a financial risk setting. It can use market or regulatory information only if that information was encoded during training or manually attached as a feature. It also gives reviewers little direct evidence for why a high-risk case was flagged. Retrieval-augmented generation (RAG) and dense retrieval methods provide a way to condition decisions on external evidence available at inference time [35,23,26,19]. In parallel, work on explanation and calibration has shown that model outputs need to be inspectable as well as accurate [34,29,14,4,15,32,40]. These observations motivate a risk model that retrieves time-valid financial evidence and links that evidence to both prediction and explanation.

This work develops the **Financial Risk Assessment Agent (FRAA)**, a retrieval-augmented model for explainable financial risk assessment. The word “agent” is used in a constrained operational sense. For each scoring request, FRAA builds a user state, queries a timestamped financial knowledge base, fuses retrieved evidence with behavioral features, estimates a risk probability, and generates a short explanation grounded in the retrieved context and influential feature groups. FRAA is therefore a structured retrieve-fuse-score-explain pipeline rather than an open-ended autonomous system.

The main contributions are:

- We design a retrieval-augmented financial risk model that combines time-aware sequence encoding, external evidence retrieval, cross-attention evidence fusion, calibrated risk scoring, and explanation generation.
- We evaluate FRAA on a large-scale proprietary multi-source dataset with a strict chronological split, and compare it with deployed tree-based baselines and neural sequence models.
- We report component ablations, feature-group occlusion, retrieval-depth analysis, explanation quality assessment, and latency measurements to examine the contribution and operational cost of the retrieval and explanation components.

## 2 Related Work

### 2.1 Financial Risk Modeling and Tabular Baselines

Credit and financial risk prediction has long relied on statistical classifiers, scorecards, tree ensembles, support vector machines, and gradient boosting. These methods handle heterogeneous tabular features and have strong empirical performance across many credit-scoring studies [16,38,3,1,28,7,13,9,25]. Prior work also shows that credit and fraud datasets are often imbalanced and operationally asymmetric, so ranking quality and calibration matter in addition to raw classification accuracy [11,8,24,5,6,31]. In production, these models are commonly supported by feature stores and domain-specific aggregation rules. Their limitation is not weak performance on static tables, but the maintenance cost of

hand-crafted features and the difficulty of using newly emerging external evidence at inference time.

## 2.2 Sequential User Modeling

Sequential recommendation and behavior modeling show that recurrent networks and Transformer encoders can capture evolving user states [17,10,39,22,36]. Financial risk assessment has a related temporal structure. Recent transactions, account changes, and market signals form irregular sequences whose relevance depends on timing and context. Sequence models have also been applied to credit-card fraud detection, where temporal order is central to identifying abnormal behavior [21]. FRAA follows this line of work but adapts it to calibrated risk scoring, where explanation is part of the target behavior rather than an optional post-processing step.

## 2.3 Retrieval-Augmented and Explainable AI

Dense retrieval and RAG methods condition model outputs on non-parametric external evidence [35,23,26,19]. Text encoders such as BERT and Sentence-BERT provide practical document representations for semantic matching [12,33]. Approximate nearest-neighbor methods and vector search libraries such as Facebook AI Similarity Search (FAISS) make such retrieval feasible at large scale [30,20]. In regulated domains, however, retrieval alone is not sufficient. Local explanation methods such as LIME and SHAP [34,29], broader XAI surveys [14,4], and calibration analyses [15,32,40] all point to the same requirement: a risk model should expose not only a probability, but also the evidence used to support it.

# 3 Method

## 3.1 Problem Formulation

For a user  $u$  and decision time  $t$ , the task is to estimate the probability  $\hat{r}_t \in [0, 1]$  that a material adverse event will occur within the next 30 days. The label  $y_t \in \{0, 1\}$  indicates whether such an event is observed. The model receives a time-ordered window of behavioral, profile, macroeconomic, and external-document signals available no later than  $t$ . This chronological constraint is applied throughout training and testing to avoid future information leakage.

Each training instance is built from an observation window ending at  $t$  and a prediction window that immediately follows it. Dynamic events after  $t$  are excluded. Slowly changing attributes are read from the latest snapshot available before  $t$ . External documents follow the same rule: a document can be retrieved only if its publication or ingestion timestamp is no later than the scoring time. This setup reflects operational risk review, where decisions are made before the adverse event and before later market information becomes available.

### 3.2 Multi-Source State Construction

FRAA maps the available evidence into three feature families. The dynamic transactional context  $\mathbf{d}_t \in \mathbb{R}^{D_d}$  summarizes recent transaction amount, frequency, volatility, merchant category, and event timing. The quasi-static profile context  $\mathbf{s}_t \in \mathbb{R}^{D_s}$  contains slowly changing user attributes, including account tenure, product holdings, balance ranges, and risk tier. The market and macro context  $\mathbf{m}_t \in \mathbb{R}^{D_m}$  contains time-aligned external indicators available before the scoring time. The external-document branch is motivated by evidence that news and regulatory language can carry risk-relevant signals [37,27].

Let  $D_d$ ,  $D_s$ , and  $D_m$  denote the dimensions of the three feature families. These vectors are concatenated and projected into a shared hidden space:

$$\mathbf{x}_t = \mathbf{W}_p[\mathbf{d}_t \parallel \mathbf{s}_t \parallel \mathbf{m}_t] + \mathbf{b}_p, \quad (1)$$

where  $\parallel$  denotes concatenation,  $\mathbf{W}_p \in \mathbb{R}^{D_h \times (D_d + D_s + D_m)}$  and  $\mathbf{b}_p \in \mathbb{R}^{D_h}$  are trainable parameters, and  $\mathbf{x}_t \in \mathbb{R}^{D_h}$  is the initial user-state representation.

### 3.3 Time-Aware Sequence Encoding

Financial behavior is observed at irregular intervals. FRAA therefore injects a time-gap encoding into each event representation. Let  $\Delta\tau_i$  denote the elapsed time between consecutive events. The input to layer  $l$  of the Transformer encoder is updated as

$$\mathbf{h}_i^{(l)} = \text{TransformerBlock}^{(l)} \left( \mathbf{h}_i^{(l-1)} + \mathbf{p}_{\Delta\tau_i} \right), \quad l = 1, \dots, L, \quad (2)$$

where  $\mathbf{p}_{\Delta\tau_i}$  is the time-gap embedding. The final state  $\mathbf{h}_t^{(L)}$  summarizes the user’s observed behavior up to time  $t$ .

### 3.4 Retrieval-Augmented Evidence Fusion

To incorporate external knowledge, FRAA converts the encoded user state into a retrieval query:

$$\mathbf{q}_t = f_q(\mathbf{h}_t^{(L)}), \quad (3)$$

where  $f_q(\cdot)$  is a lightweight multilayer perceptron that maps the sequence state to the document-embedding space. The financial knowledge base contains approximately 1.2 million timestamped documents, including regulatory filings, market reports, news articles, and macroeconomic commentaries. Documents are embedded offline and indexed with FAISS. At inference time, FRAA retrieves the top- $K$  documents by cosine similarity:

$$\mathcal{D}_t = \text{topK}_{d \in \mathcal{K}_t} \frac{\mathbf{q}_t^\top \mathbf{e}_d}{\|\mathbf{q}_t\| \|\mathbf{e}_d\|}, \quad (4)$$

where  $\mathcal{K}_t$  contains only documents available no later than  $t$ , and  $\mathbf{e}_d$  is the embedding of document  $d$ . Document embeddings are produced by a financial-domain

text encoder initialized from bidirectional Transformer encoders and sentence-level retrieval models [12,33]. The timestamp restriction prevents future documents from leaking into the prediction.

The retrieved document embeddings are fused with the user state using cross-attention:

$$\tilde{\mathbf{h}}_t = \text{CrossAttn} \left( \mathbf{h}_t^{(L)} \mathbf{W}_Q, \mathbf{D}_t \mathbf{W}_K, \mathbf{D}_t \mathbf{W}_V \right), \quad (5)$$

where  $\mathbf{D}_t \in \mathbb{R}^{K \times D_e}$  stacks the embeddings of retrieved documents,  $D_e$  is the document-embedding dimension, and  $\mathbf{W}_Q$ ,  $\mathbf{W}_K$ , and  $\mathbf{W}_V$  are trainable projection matrices. The fused representation  $\tilde{\mathbf{h}}_t$  is passed to the scoring and explanation heads. Cross-attention is used instead of simple concatenation because the relevance of external evidence depends on the user state. For example, a macroeconomic commentary may matter for users with concentrated sector exposure, but may be weakly relevant for users whose recent transactions do not involve that sector.

### 3.5 Risk Scoring and Explanation Generation

The risk head maps the fused representation to a calibrated probability:

$$\hat{r}_t = \sigma(\mathbf{w}_r^\top \tilde{\mathbf{h}}_t + b_r). \quad (6)$$

where  $\sigma(\cdot)$  is the sigmoid function,  $\mathbf{w}_r$  is a trainable risk-head vector, and  $b_r$  is a scalar bias. In parallel, the explanation head generates a concise textual justification from the retrieved documents and feature-group importance signals. The training objective combines per-instance binary cross-entropy for risk prediction with a supervised explanation-fidelity term:

$$\mathcal{L} = -y_t \log \hat{r}_t - (1 - y_t) \log(1 - \hat{r}_t) + \lambda \mathcal{L}_{\text{exp}}(\hat{e}_t, e_t^{\text{ref}}), \quad (7)$$


where  $\hat{e}_t$  is the generated explanation,  $e_t^{\text{ref}}$  is an expert or rule-supported reference explanation,  $\mathcal{L}_{\text{exp}}$  penalizes mismatches with reference evidence and feature-group rationales, and  $\lambda$  balances prediction and explanation supervision. The explanation head is trained to refer to retrieved evidence and influential feature groups, rather than to produce a free-form rationale disconnected from the risk score. In all reported experiments, the Transformer uses 8 attention blocks, hidden dimension 256, and embedding dimension 512. Retrieval depth is  $K = 5$  unless otherwise stated.

Figure 1 summarizes the inference path from multi-source evidence encoding to retrieval, fusion, scoring, and explanation generation.

## 4 Experiments

### 4.1 Dataset and Evaluation Protocol

The dataset is a proprietary 24-month multi-source corpus from a financial risk environment. It contains over 20 million users, 1.2 billion interactions, and 15



**Fig. 1.** Overview of FRAA. The model encodes multi-source financial behavior, retrieves timestamp-valid documents from a financial knowledge base, fuses retrieved evidence with the user state, and outputs both a calibrated risk score and an explanation.

heterogeneous data sources. Records include transaction events, account status snapshots, product holdings, macroeconomic series, and timestamped news or regulatory signals. The target indicates whether a material adverse event, such as default, severe delinquency, or forced liquidation, occurs within the next 30 days.

We use a strict chronological split. Historical records are used to build look-back features and timestamp-valid knowledge-base entries. The model is trained on a 12-month period, tuned on the following month, and evaluated on the final month. No validation or test-period documents are made available to earlier predictions. We report Log Loss for probability calibration, Recall@1/Recall@5 for risk-review ranking, area under the receiver operating characteristic curve (AUC) for feature-group occlusion, and longest common subsequence ROUGE (ROUGE-L) for explanation overlap with expert-written references.

The evaluation protocol follows the main failure modes of financial risk systems. Log Loss measures whether predicted probabilities are useful for downstream thresholding and risk appetite control. Recall@1 and Recall@5 measure the ranking quality of the highest-risk cases, which are the cases most likely to be reviewed by analysts. The chronological split and timestamp-valid retrieval rule are used for every model variant, including ablations, so that performance differences do not arise from different access to future information.

## 4.2 Baselines and Training Details

We compare FRAA with a deployed Random Forest (RF), RF with static knowledge-base (KB) features, XGBoost, LSTM, and a non-retrieval Transformer. These baselines cover the main operational choices in the target setting:

tabular ensembles, tabular ensembles with manually attached external features, recurrent sequence modeling, and neural sequence modeling without retrieval. We retain the deployed tree-based baseline because prior credit-scoring and tabular-learning studies repeatedly find tree ensembles difficult to displace on heterogeneous tabular data [3,8,2]. All baselines use the same chronological split and available feature families. FRAA is implemented in PyTorch and trained with AdamW, learning rate  $1 \times 10^{-4}$ , weight decay  $1 \times 10^{-5}$ , batch size 64, and a maximum budget of 50 epochs. The selected checkpoint is the one with the lowest validation Log Loss.

Hyperparameters are chosen on the validation month and then fixed for the test month. For retrieval variants, document embeddings are constructed offline and indexed before evaluation. During scoring, the retrieval pool is filtered by timestamp before nearest-neighbor search. This step is essential because a high-capacity retrieval model could otherwise use documents that were unavailable when the decision should have been made.

### 4.3 Evaluation Questions

The experiments address three questions. First, does retrieval-augmented evidence change calibrated risk prediction and high-risk ranking relative to tree-based and neural baselines? Second, which components and evidence sources account for the observed gains, especially the dynamic transaction encoder, time-gap encoding, and knowledge retrieval module? Third, are the explanations and latency compatible with analyst-facing risk review? The following sections follow this order, separating the main predictive result from auxiliary analyses.

### 4.4 Main Results

We first compare the full model with deployed and neural baselines under the same chronological split. Table 1 reports the numerical results, and Fig. 2 visualizes the two primary metrics. The ordering is consistent across calibration and ranking: the full retrieval-augmented model outperforms tree-based baselines, the non-retrieval Transformer, and all ablated FRAA variants.

Relative to the deployed RF baseline, FRAA reduces Log Loss by 19.6% and raises Recall@5 from 0.3172 to 0.4185. Relative to the non-retrieval Transformer, it reduces Log Loss by 9.4% and raises Recall@5 by 8.8%. The ablated variants remain below the full model, indicating that retrieval, temporal encoding, and dynamic features contribute complementary information.


## 5 Ablation and Analysis

### 5.1 Architectural Components

The component ablation removes one design choice at a time while keeping the training split, feature availability, and evaluation metrics unchanged. This setup separates the effect of retrieval from the effect of stronger sequence modeling.

**Table 1.** Main test-set performance. Lower Log Loss and higher Recall@5 are better.

| Model                  | Log Loss      | Recall@5      |
|------------------------|---------------|---------------|
| RF (deployed baseline) | 0.6421        | 0.3172        |
| RF + KB features       | 0.6289        | 0.3291        |
| XGBoost                | 0.6104        | 0.3495        |
| LSTM                   | 0.5873        | 0.3738        |
| Transformer            | 0.5698        | 0.3846        |
| FRAA (-Dyn)            | 0.5642        | 0.3891        |
| FRAA (-TE)             | 0.5510        | 0.3957        |
| FRAA (-KB)             | 0.5385        | 0.4023        |
| <b>FRAA (full)</b>     | <b>0.5163</b> | <b>0.4185</b> |

**Fig. 2.** Main-result comparison. FRAA obtains the lowest Log Loss and highest Recall@5 among all evaluated models.**Table 2.** Component ablation on the risk-score prediction task.


| Model variant      | Log Loss      | Recall@5      |
|--------------------|---------------|---------------|
| Transformer        | 0.5698        | 0.3846        |
| FRAA (-Dyn)        | 0.5642        | 0.3891        |
| FRAA (-TE)         | 0.5510        | 0.3957        |
| FRAA (-KB)         | 0.5385        | 0.4023        |
| <b>FRAA (full)</b> | <b>0.5163</b> | <b>0.4185</b> |



Table 2 isolates the contribution of dynamic transactions (-Dyn), temporal encoding (-TE), and knowledge retrieval (-KB). All three removals degrade the full model, confirming that the gains do not come from a single add-on. Removing the knowledge component increases Log Loss from 0.5163 to 0.5385 and decreases Recall@5 from 0.4185 to 0.4023, supporting the retrieval-augmented design. Removing temporal encoding also reduces performance, which is consistent with the irregular timing of financial events. The -Dyn variant remains stronger than the non-retrieval Transformer. External evidence and temporal encoding therefore compensate for some missing dynamic transaction information, but they do not replace it.

## 5.2 Feature Groups and Retrieval Depth

We next examine the gain from two complementary views: feature-family occlusion and retrieval depth. Figure 3 summarizes both analyses using the same values as the feature-family occlusion and retrieval-depth experiments. Knowledge-retrieved context produces the largest AUC drop, followed by transaction dynamics. Static profile variables still contribute, but they are less influential than time-sensitive behavioral and external-evidence signals. Retrieval depth also has a clear optimum in this setting: one or three documents provide insufficient external context, whereas seven or ten documents slightly degrade both metrics. We therefore use  $K = 5$  in the main experiments. The degradation at larger  $K$  is small, but it is consistent with the retrieval-noise trade-off: broader retrieval increases coverage while introducing less relevant documents into the fusion module.




**Fig. 3.** Retrieval evidence analysis. (a) Feature-family occlusion measured by AUC drop. (b) Retrieval-depth sensitivity, with  $K = 5$  selected for the main experiments.

## 6 Further Evaluation

### 6.1 Explanation Quality

Beyond risk scores, FRAA is intended to support human review by explaining why a high-risk case was flagged. We therefore evaluate the generated explanations using both expert ratings and lexical overlap with reference explanations.



**Fig. 4.** Explanation-quality evaluation. (a) Expert-rated faithfulness and readability on a five-point scale. (b) ROUGE-L overlap with reference explanations.

Five domain experts evaluated 1,000 high-risk test predictions. Figure 4 reports higher faithfulness, readability, and ROUGE-L for FRAA than for a template baseline that lists the top three feature groups. The full model obtains faithfulness of 4.32/5, readability of 4.51/5, and ROUGE-L of 0.4131. The result suggests that the explanation head adds information beyond feature enumeration. The expert ratings should be interpreted as an auditability signal, not as evidence that generated text is sufficient for automatic decision making.


### 6.2 Efficiency Analysis

The final evaluation asks whether the added retrieval and fusion steps are practical for an analyst-facing review loop.

FRAA is slower than tree ensembles because it performs retrieval and neural sequence encoding. Even so, the tested latency in Fig. 5 remains compatible with interactive risk review: FRAA requires 15.7 ms per query and reaches 4,070 queries/s in the tested serving configuration. The retrieval step accounts for approximately 4.2 ms of the total latency. Throughput is lower than for tabular baselines, but remains above the level required for batched offline scoring and analyst triage in the tested configuration.

## 7 Limitations and Reproducibility

The evaluation uses a proprietary financial dataset that cannot be publicly released because it contains sensitive user-level risk and transaction information.



**Fig. 5.** Inference latency and throughput under the tested batch-size-64 serving configuration.

This limits independent replication. To support reproducibility within this constraint, we report the data scale, feature families, chronological split, training configuration, retrieval depth, baselines, and ablation settings. The knowledge base is timestamped, and retrieval uses only documents available at or before the scoring time. These details make the experimental protocol auditable even when the raw records cannot be shared.

Several claims should be read within this scope. First, the results establish performance in one large multi-source financial environment, not universal transfer across institutions. Second, the current evaluation does not report multi-seed variance or formal statistical significance tests because the available audit record contains one fixed proprietary chronological split rather than repeated independent training runs. Third, the explanation evaluation measures expert-perceived faithfulness and readability for high-risk test predictions. It does not remove the need for human oversight in regulated decisions. Future work should add public benchmark evaluations where legally possible, repeatability analysis under multiple splits or random seeds, institution-level transfer tests including parameter-efficient adaptation methods such as LoRA [18], and governance mechanisms for retrieved evidence that may become outdated or disputed.

## 8 Conclusion

This paper presented FRAA, a retrieval-augmented model for explainable financial risk assessment. FRAA combines time-aware sequence modeling, timestamp-valid knowledge retrieval, cross-attention evidence fusion, calibrated risk scoring, and explanation generation. On a large-scale proprietary multi-source dataset, FRAA achieves the lowest Log Loss and highest Recall@5 among the evaluated models. Ablations show that dynamic behavior, temporal encoding, and retrieval provide complementary evidence, while feature-group occlusion identi-

fies knowledge-retrieved context as the strongest feature-family signal. Retrieval-depth analysis selects five documents as the best trade-off between evidence coverage and noise. Additional evaluations examine explanation quality and tested latency. Taken together, the results suggest that external evidence retrieval is most valuable when it is constrained by time validity, fused with temporal user behavior, and exposed through auditable explanations.

## References

1. Abdou, H.A., Pointon, J.: Credit scoring, statistical techniques and evaluation criteria: A review of the literature. *Intelligent Systems in Accounting, Finance and Management* **18**(2-3), 59–88 (2011). <https://doi.org/10.1002/isaf.325>
2. Arik, S.O., Pfister, T.: Tabnet: Attentive interpretable tabular learning. *Proceedings of the AAAI Conference on Artificial Intelligence* **35**(8), 6679–6687 (2021). <https://doi.org/10.1609/aaai.v35i8.16826>
3. Baesens, B., Van Gestel, T., Viaene, S., Stepanova, M., Suykens, J., Vanthienen, J.: Benchmarking state-of-the-art classification algorithms for credit scoring. *Journal of the Operational Research Society* **54**(6), 627–635 (2003). <https://doi.org/10.1057/palgrave.jors.2601545>
4. Barredo Arrieta, A., Diaz-Rodriguez, N., Del Ser, J., Benetot, A., Tabik, S., Barabado, A., Garcia, S., Gil-Lopez, S., Molina, D., Benjamins, R., Chatila, R., Herrera, F.: Explainable artificial intelligence (xai): Concepts, taxonomies, opportunities and challenges toward responsible ai. *Information Fusion* **58**, 82–115 (2020). <https://doi.org/10.1016/j.inffus.2019.12.012>
5. Bellotti, T., Crook, J.: Support vector machines for credit scoring and discovery of significant features. *Expert Systems with Applications* **36**(2), 3302–3308 (2009). <https://doi.org/10.1016/j.eswa.2008.01.005>
6. Bolton, R.J., Hand, D.J.: Statistical fraud detection: A review. *Statistical Science* **17**(3), 235–255 (2002). <https://doi.org/10.1214/ss/1042727940>
7. Breiman, L.: Random forests. *Machine Learning* **45**(1), 5–32 (2001). <https://doi.org/10.1023/A:1010933404324>
8. Brown, I., Mues, C.: An experimental comparison of classification algorithms for imbalanced credit scoring data sets. *Expert Systems with Applications* **39**(3), 3446–3453 (2012). <https://doi.org/10.1016/j.eswa.2011.09.033>
9. Chen, T., Guestrin, C.: Xgboost: A scalable tree boosting system. In: *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. pp. 785–794 (2016). <https://doi.org/10.1145/2939672.2939785>
10. Cho, K., van Merriënboer, B., Gulcehre, C., Bahdanau, D., Bougares, F., Schwenk, H., Bengio, Y.: Learning phrase representations using rnn encoder-decoder for statistical machine translation. In: *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing*. pp. 1724–1734 (2014). <https://doi.org/10.3115/v1/D14-1179>
11. Crook, J.N., Edelman, D.B., Thomas, L.C.: Recent developments in consumer credit risk assessment. *European Journal of Operational Research* **183**(3), 1447–1465 (2007). <https://doi.org/10.1016/j.ejor.2006.09.100>
12. Devlin, J., Chang, M.W., Lee, K., Toutanova, K.: Bert: Pre-training of deep bidirectional transformers for language understanding. In: *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*. pp. 4171–4186 (2019). <https://doi.org/10.18653/v1/N19-1423>

13. Friedman, J.H.: Greedy function approximation: A gradient boosting machine. *The Annals of Statistics* **29**(5), 1189–1232 (2001). <https://doi.org/10.1214/aos/1013203451>
14. Guidotti, R., Monreale, A., Ruggieri, S., Turini, F., Giannotti, F., Pedreschi, D.: A survey of methods for explaining black box models. *ACM Computing Surveys* **51**(5), 1–42 (2018). <https://doi.org/10.1145/3236009>
15. Guo, C., Pleiss, G., Sun, Y., Weinberger, K.Q.: On calibration of modern neural networks. In: *Proceedings of the 34th International Conference on Machine Learning*. pp. 1321–1330 (2017)
16. Hand, D.J., Henley, W.E.: Statistical classification methods in consumer credit scoring: A review. *Journal of the Royal Statistical Society: Series A (Statistics in Society)* **160**(3), 523–541 (1997). <https://doi.org/10.1111/j.1467-985X.1997.00078.x>
17. Hochreiter, S., Schmidhuber, J.: Long short-term memory. *Neural Computation* **9**(8), 1735–1780 (1997). <https://doi.org/10.1162/neco.1997.9.8.1735>
18. Hu, E.J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., Wang, L., Chen, W.: Lora: Low-rank adaptation of large language models. In: *Proceedings of the International Conference on Learning Representations* (2022)
19. Izacard, G., Grave, E.: Leveraging passage retrieval with generative models for open domain question answering. In: *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics*. pp. 874–880 (2021). <https://doi.org/10.18653/v1/2021.eacl-main.74>
20. Johnson, J., Douze, M., Jegou, H.: Billion-scale similarity search with gpus. *IEEE Transactions on Big Data* **7**(3), 535–547 (2021). <https://doi.org/10.1109/TBDATA.2019.2921572>
21. Jurgovsky, J., Granitzer, M., Ziegler, K., Calabretto, S., Portier, P.E., He-Guelton, L., Caelen, O.: Sequence classification for credit-card fraud detection. *Expert Systems with Applications* **100**, 234–245 (2018). <https://doi.org/10.1016/j.eswa.2018.01.037>
22. Kang, W.C., McAuley, J.: Self-attentive sequential recommendation. In: *Proceedings of the IEEE International Conference on Data Mining*. pp. 197–206 (2018). <https://doi.org/10.1109/ICDM.2018.00035>
23. Karpukhin, V., Oguz, B., Min, S., Lewis, P., Wu, L., Edunov, S., Chen, D., Yih, W.t.: Dense passage retrieval for open-domain question answering. In: *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*. pp. 6769–6781 (2020). <https://doi.org/10.18653/v1/2020.emnlp-main.550>
24. Khandani, A.E., Kim, A.J., Lo, A.W.: Consumer credit-risk models via machine-learning algorithms. *Journal of Banking and Finance* **34**(11), 2767–2787 (2010). <https://doi.org/10.1016/j.jbankfin.2010.06.001>
25. Lessmann, S., Baesens, B., Seow, H.V., Thomas, L.C.: Benchmarking state-of-the-art classification algorithms for credit scoring: An update of research. *European Journal of Operational Research* **247**(1), 124–136 (2015). <https://doi.org/10.1016/j.ejor.2015.05.030>
26. Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Kuttler, H., Lewis, M., Yih, W.t., Rocktaschel, T., Riedel, S., Kiela, D.: Retrieval-augmented generation for knowledge-intensive nlp tasks. In: *Advances in Neural Information Processing Systems*. vol. 33, pp. 9459–9474 (2020)
27. Loughran, T., McDonald, B.: When is a liability not a liability? textual analysis, dictionaries, and 10-ks. *The Journal of Finance* **66**(1), 35–65 (2011). <https://doi.org/10.1111/j.1540-6261.2010.01625.x>

28. Louzada, F., Ara, A., Fernandes, G.B.: Classification methods applied to credit scoring: Systematic review and overall comparison. *Surveys in Operations Research and Management Science* **21**(2), 117–134 (2016). <https://doi.org/10.1016/j.sorms.2016.10.001>
29. Lundberg, S.M., Lee, S.I.: A unified approach to interpreting model predictions. In: *Advances in Neural Information Processing Systems*. pp. 4765–4774 (2017)
30. Malkov, Y.A., Yashunin, D.A.: Efficient and robust approximate nearest neighbor search using hierarchical navigable small world graphs. *IEEE Transactions on Pattern Analysis and Machine Intelligence* **42**(4), 824–836 (2020). <https://doi.org/10.1109/TPAMI.2018.2889473>
31. Ngai, E.W.T., Hu, Y., Wong, Y.H., Chen, Y., Sun, X.: The application of data mining techniques in financial fraud detection: A classification framework and an academic review of literature. *Decision Support Systems* **50**(3), 559–569 (2011). <https://doi.org/10.1016/j.dss.2010.08.006>
32. Niculescu-Mizil, A., Caruana, R.: Predicting good probabilities with supervised learning. In: *Proceedings of the 22nd International Conference on Machine Learning*. pp. 625–632 (2005). <https://doi.org/10.1145/1102351.1102430>
33. Reimers, N., Gurevych, I.: Sentence embeddings using siamese bert-networks. In: *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing*. pp. 3980–3990 (2019). <https://doi.org/10.18653/v1/D19-1410>
34. Ribeiro, M.T., Singh, S., Guestrin, C.: “why should i trust you?”: Explaining the predictions of any classifier. In: *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. pp. 1135–1144 (2016). <https://doi.org/10.1145/2939672.2939778>
35. Robertson, S., Zaragoza, H.: The probabilistic relevance framework: Bm25 and beyond. *Foundations and Trends in Information Retrieval* **3**(4), 333–389 (2009). <https://doi.org/10.1561/15000000019>
36. Sun, F., Liu, J., Wu, J., Pei, C., Lin, X., Ou, W., Jiang, P.: Bert4rec: Sequential recommendation with bidirectional encoder representations from transformer. In: *Proceedings of the 28th ACM International Conference on Information and Knowledge Management*. pp. 1441–1450 (2019). <https://doi.org/10.1145/3357384.3357895>
37. Tetlock, P.C.: Giving content to investor sentiment: The role of media in the stock market. *The Journal of Finance* **62**(3), 1139–1168 (2007). <https://doi.org/10.1111/j.1540-6261.2007.01232.x>
38. Thomas, L.C.: A survey of credit and behavioural scoring: Forecasting financial risk of lending to consumers. *International Journal of Forecasting* **16**(2), 149–172 (2000). [https://doi.org/10.1016/S0169-2070\(00\)00034-0](https://doi.org/10.1016/S0169-2070(00)00034-0)
39. Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A.N., Kaiser, L., Polosukhin, I.: Attention is all you need. In: *Advances in Neural Information Processing Systems*. pp. 5998–6008 (2017)
40. Zadrozny, B., Elkan, C.: Transforming classifier scores into accurate multiclass probability estimates. In: *Proceedings of the Eighth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. pp. 694–699 (2002). <https://doi.org/10.1145/775047.775151>